

Data Mining

Chapter 5

Association Analysis: Basic Concepts

Introduction to Data Mining, 2nd Edition

by

Tan, Steinbach, Karpatne, Kumar

3/8/2021

Introduction to Data Mining, 2nd Edition

1

1

Association Rule Mining

- Given a set of transactions, find rules that will predict the occurrence of an item based on the occurrences of other items in the transaction

Market-Basket transactions

TID	Items
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

Example of Association Rules

$\{\text{Diaper}\} \rightarrow \{\text{Beer}\},$
 $\{\text{Milk, Bread}\} \rightarrow \{\text{Eggs, Coke}\},$
 $\{\text{Beer, Bread}\} \rightarrow \{\text{Milk}\},$

Implication means co-occurrence,
not causality!

3/8/2021

Introduction to Data Mining, 2nd Edition

2

2

Definition: Frequent Itemset

- **Itemset**

- A collection of one or more items
 - ◆ Example: {Milk, Bread, Diaper}
- k-itemset
 - ◆ An itemset that contains k items

- **Support count (σ)**

- Frequency of occurrence of an itemset
- E.g. $\sigma(\{\text{Milk, Bread, Diaper}\}) = 2$

- **Support**

- Fraction of transactions that contain an itemset
- E.g. $s(\{\text{Milk, Bread, Diaper}\}) = 2/5$

- **Frequent Itemset**

- An itemset whose support is greater than or equal to a *minsup* threshold

TID	Items
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

3/8/2021

Introduction to Data Mining, 2nd Edition

3

3

Definition: Association Rule

- **Association Rule**

- An implication expression of the form $X \rightarrow Y$, where X and Y are itemsets
- Example:
 $\{\text{Milk, Diaper}\} \rightarrow \{\text{Beer}\}$

- **Rule Evaluation Metrics**

- Support (s)
 - ◆ Fraction of transactions that contain both X and Y
- Confidence (c)
 - ◆ Measures how often items in Y appear in transactions that contain X

Example:

$\{\text{Milk, Diaper}\} \Rightarrow \{\text{Beer}\}$

$$s = \frac{\sigma(\text{Milk, Diaper, Beer})}{|T|} = \frac{2}{5} = 0.4$$

$$c = \frac{\sigma(\text{Milk, Diaper, Beer})}{\sigma(\text{Milk, Diaper})} = \frac{2}{3} = 0.67$$

TID	Items
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

3/8/2021

Introduction to Data Mining, 2nd Edition

4

4

Association Rule Mining Task

- Given a set of transactions T , the goal of association rule mining is to find all rules having
 - support $\geq \text{minsup}$ threshold
 - confidence $\geq \text{minconf}$ threshold
- Brute-force approach:
 - List all possible association rules
 - Compute the support and confidence for each rule
 - Prune rules that fail the *minsup* and *minconf* thresholds

⇒ **Computationally prohibitive!**

3/8/2021


Introduction to Data Mining, 2nd Edition

5

5

Computational Complexity

- Given d unique items:
 - Total number of itemsets $= 2^d$
 - Total number of possible association rules:



$$R = \sum_{k=1}^{d-1} \left[\binom{d}{k} \times \sum_{j=1}^{d-k} \binom{d-k}{j} \right]$$
$$= 3^d - 2^{d+1} + 1$$

If $d=6$, $R = 602$ rules

3/8/2021

Introduction to Data Mining, 2nd Edition

6

6

Mining Association Rules

TID	Items
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

Example of Rules:

$\{\text{Milk, Diaper}\} \rightarrow \{\text{Beer}\}$ ($s=0.4, c=0.67$)
 $\{\text{Milk, Beer}\} \rightarrow \{\text{Diaper}\}$ ($s=0.4, c=1.0$)
 $\{\text{Diaper, Beer}\} \rightarrow \{\text{Milk}\}$ ($s=0.4, c=0.67$)
 $\{\text{Beer}\} \rightarrow \{\text{Milk, Diaper}\}$ ($s=0.4, c=0.67$)
 $\{\text{Diaper}\} \rightarrow \{\text{Milk, Beer}\}$ ($s=0.4, c=0.5$)
 $\{\text{Milk}\} \rightarrow \{\text{Diaper, Beer}\}$ ($s=0.4, c=0.5$)

Observations:

- All the above rules are binary partitions of the same itemset:
 $\{\text{Milk, Diaper, Beer}\}$
- Rules originating from the same itemset have identical support but can have different confidence
- Thus, we may decouple the support and confidence requirements

3/8/2021

Introduction to Data Mining, 2nd Edition

7

7

Mining Association Rules

- Two-step approach:
 1. **Frequent Itemset Generation**
 - Generate all itemsets whose support \geq minsup
 2. **Rule Generation**
 - Generate high confidence rules from each frequent itemset, where each rule is a binary partitioning of a frequent itemset
- Frequent itemset generation is still computationally expensive

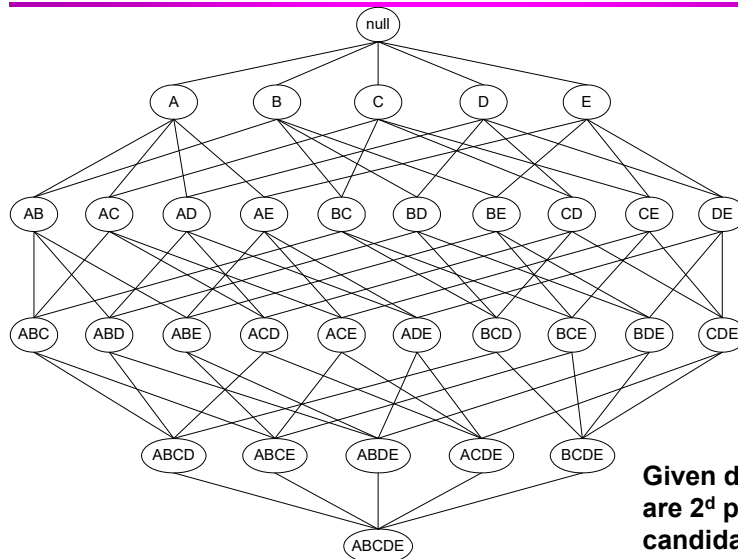
3/8/2021

Introduction to Data Mining, 2nd Edition

8

8

Frequent Itemset Generation



Given d items, there are 2^d possible candidate itemsets

3/8/2021

Introduction to Data Mining, 2nd Edition

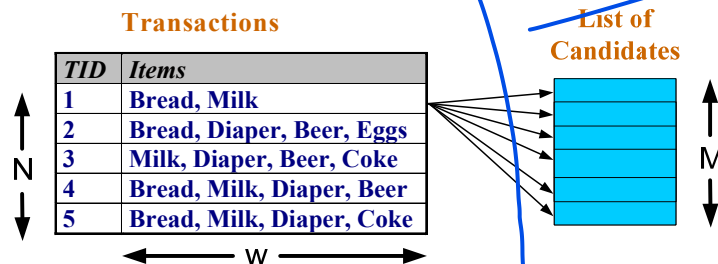
9

9

Frequent Itemset Generation

● Brute-force approach:

- Each itemset in the lattice is a **candidate** frequent itemset
- Count the support of each candidate by scanning the database



- Match each transaction against every candidate
- Complexity $\sim O(NMw) \Rightarrow$ **Expensive since $M = 2^d$!!!**

3/8/2021

Introduction to Data Mining, 2nd Edition

10

10

Frequent Itemset Generation Strategies

- Reduce the **number of candidates** (M)
 - Complete search: $M=2^d$
 - Use pruning techniques to reduce M
- Reduce the **number of transactions** (N)
 - Reduce size of N as the size of itemset increases
 - Used by DHP and vertical-based mining algorithms
- Reduce the **number of comparisons** (NM)
 - Use efficient data structures to store the candidates or transactions
 - No need to match every candidate against every transaction

3/8/2021

Introduction to Data Mining, 2nd Edition

11

11

Reducing Number of Candidates

- **Apriori principle:**
 - If an itemset is frequent, then all of its subsets must also be frequent
- Apriori principle holds due to the following property of the support measure:

$$\forall X, Y : (X \subseteq Y) \Rightarrow s(X) \geq s(Y)$$

- Support of an itemset never exceeds the support of its subsets
- This is known as the **anti-monotone** property of support

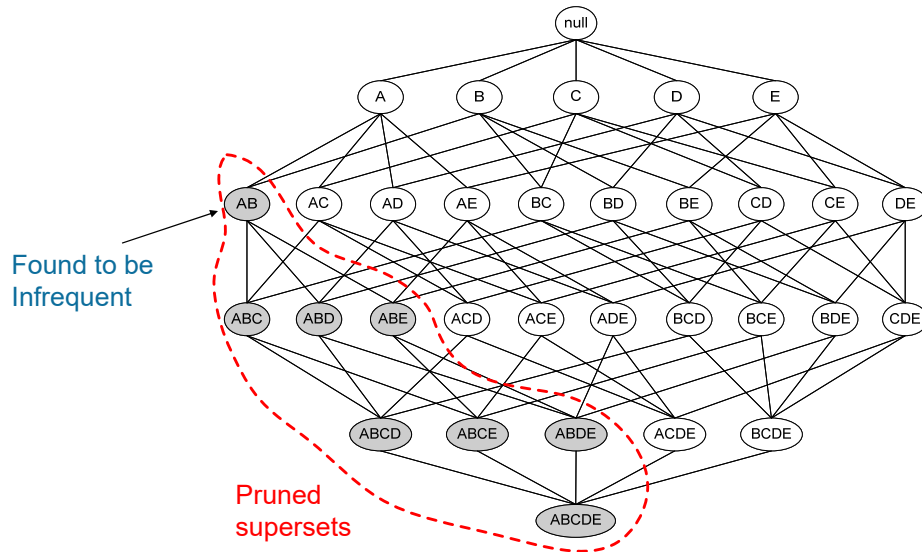
3/8/2021

Introduction to Data Mining, 2nd Edition

12

12

Illustrating Apriori Principle



3/8/2021

Introduction to Data Mining, 2nd Edition

13

13

Illustrating Apriori Principle

TID	Items
1	Bread, Milk
2	Beer, Bread, Diaper, Eggs
3	Beer, Coke, Diaper, Milk
4	Beer, Bread, Diaper, Milk
5	Bread, Coke, Diaper, Milk

Items (1-itemsets)

Item	Count
Bread	4
Coke	2
Milk	4
Beer	3
Diaper	4
Eggs	1

Minimum Support = 3

If every subset is considered,
 ${}^6C_1 + {}^6C_2 + {}^6C_3$
 $6 + 15 + 20 = 41$
 With support-based pruning,
 $6 + 6 + 4 = 16$

3/8/2021

Introduction to Data Mining, 2nd Edition

14

14

Illustrating Apriori Principle

TID	Items
1	Bread, Milk
2	Beer, Bread, Diaper, Eggs
3	Beer, Coke, Diaper, Milk
4	Beer, Bread, Diaper, Milk
5	Bread, Coke, Diaper, Milk



Items (1-itemsets)

Item	Count
Bread	4
Coke	2
Milk	4
Beer	3
Diaper	4
Eggs	1

Minimum Support = 3

If every subset is considered,
 ${}^6C_1 + {}^6C_2 + {}^6C_3$
 $6 + 15 + 20 = 41$
 With support-based pruning,
 $6 + 6 + 4 = 16$

15

Illustrating Apriori Principle

TID	Items
1	Bread, Milk
2	Beer, Bread, Diaper, Eggs
3	Beer, Coke, Diaper, Milk
4	Beer, Bread, Diaper, Milk
5	Bread, Coke, Diaper, Milk

Item	Count
Bread	4
Coke	2
Milk	4
Beer	3
Diaper	4
Eggs	1

Items (1-itemsets)



Itemset
{Bread, Milk}
{Bread, Beer }
{Bread, Diaper}
{Beer, Milk}
{Diaper, Milk}
{Beer, Diaper}

Pairs (2-itemsets)

(No need to generate candidates involving Coke or Eggs)

Minimum Support = 3

If every subset is considered,
 ${}^6C_1 + {}^6C_2 + {}^6C_3$
 $6 + 15 + 20 = 41$
 With support-based pruning,
 $6 + 6 + 4 = 16$

16

Illustrating Apriori Principle

TID	Items
1	Bread, Milk
2	Beer, Bread, Diaper, Eggs
3	Beer, Coke, Diaper, Milk
4	Beer, Bread, Diaper, Milk
5	Bread, Coke, Diaper, Milk

Item	Count
Bread	4
Coke	2
Milk	4
Beer	3
Diaper	4
Eggs	1

Items (1-itemsets)



Itemset	Count
{Bread, Milk}	3
{Beer, Bread}	2
{Bread, Diaper}	3
{Beer, Milk}	2
{Diaper, Milk}	3
{Beer, Diaper}	3

Pairs (2-itemsets)

(No need to generate candidates involving Coke or Eggs)

Minimum Support = 3

If every subset is considered,
 ${}^6C_1 + {}^6C_2 + {}^6C_3$
 $6 + 15 + 20 = 41$
 With support-based pruning,
 $6 + 6 + 4 = 16$

3/8/2021

Introduction to Data Mining, 2nd Edition

17

17

Illustrating Apriori Principle

TID	Items
1	Bread, Milk
2	Beer, Bread, Diaper, Eggs
3	Beer, Coke, Diaper, Milk
4	Beer, Bread, Diaper, Milk
5	Bread, Coke, Diaper, Milk

Item	Count
Bread	4
Coke	2
Milk	4
Beer	3
Diaper	4
Eggs	1

Items (1-itemsets)



Itemset	Count
{Bread, Milk}	3
{Bread, Beer}	2
{Bread, Diaper}	3
{Milk, Beer}	2
{Milk, Diaper}	3
{Beer, Diaper}	3

Pairs (2-itemsets)

(No need to generate candidates involving Coke or Eggs)

Minimum Support = 3

If every subset is considered,
 ${}^6C_1 + {}^6C_2 + {}^6C_3$
 $6 + 15 + 20 = 41$
 With support-based pruning,
 $6 + 6 + 4 = 16$



Itemset
{Beer, Diaper, Milk}
{Beer, Bread, Diaper}
{Bread, Diaper, Milk}
{Beer, Bread, Milk}

Triplets (3-itemsets)

3/8/2021

Introduction to Data Mining, 2nd Edition

18

18

Illustrating Apriori Principle

TID	Items
1	Bread, Milk
2	Beer, Bread, Diaper, Eggs
3	Beer, Coke, Diaper, Milk
4	Beer, Bread, Diaper, Milk
5	Bread, Coke, Diaper, Milk

Item	Count
Bread	4
Coke	2
Milk	4
Beer	3
Diaper	4
Eggs	1

Items (1-itemsets)



Itemset	Count
{Bread, Milk}	3
{Bread, Beer}	2
{Bread, Diaper}	3
{Milk, Beer}	2
{Milk, Diaper}	3
{Beer, Diaper}	3

Pairs (2-itemsets)

(No need to generate candidates involving Coke or Eggs)



Itemset	Count
{Beer, Diaper, Milk}	2
{Beer, Bread, Diaper}	2
{Bread, Diaper, Milk}	2
{Beer, Bread, Milk}	1

Triplets (3-itemsets)

Minimum Support = 3

If every subset is considered,
 ${}^6C_1 + {}^6C_2 + {}^6C_3$
 $6 + 15 + 20 = 41$
 With support-based pruning,
 $6 + 6 + 4 = 16$

3/8/2021

Introduction to Data Mining, 2nd Edition

19

19

Illustrating Apriori Principle

TID	Items
1	Bread, Milk
2	Beer, Bread, Diaper, Eggs
3	Beer, Coke, Diaper, Milk
4	Beer, Bread, Diaper, Milk
5	Bread, Coke, Diaper, Milk

Item	Count
Bread	4
Coke	2
Milk	4
Beer	3
Diaper	4
Eggs	1

Items (1-itemsets)



Itemset	Count
{Bread, Milk}	3
{Bread, Beer}	2
{Bread, Diaper}	3
{Milk, Beer}	2
{Milk, Diaper}	3
{Beer, Diaper}	3

Pairs (2-itemsets)

(No need to generate candidates involving Coke or Eggs)



Itemset	Count
{Beer, Diaper, Milk}	2
{Beer, Bread, Diaper}	2
{Bread, Diaper, Milk}	2
{Beer, Bread, Milk}	1

Triplets (3-itemsets)

Minimum Support = 3

If every subset is considered,
 ${}^6C_1 + {}^6C_2 + {}^6C_3$
 $6 + 15 + 20 = 41$
 With support-based pruning,
 $6 + 6 + 4 = 16$
 $6 + 6 + 1 = 13$

3/8/2021

Introduction to Data Mining, 2nd Edition

20

20
 Prune ke phle w/o duplicates
 → Itno ka support nikalna pada
 Jo candidates prev stage me

prune hue UNKO Pchle h; Prune
Karo

Apriori Algorithm

- F_k : frequent k-itemsets
- L_k : candidate k-itemsets

Algorithm

- Let $k=1$
- Generate $F_1 = \{\text{frequent 1-itemsets}\}$
- Repeat until F_k is empty
 - ◆ **Candidate Generation**: Generate L_{k+1} from F_k
 - ◆ **Candidate Pruning**: Prune candidate itemsets in L_{k+1} containing subsets of length k that are infrequent
 - ◆ **Support Counting**: Count the support of each candidate in L_{k+1} by scanning the DB
 - ◆ **Candidate Elimination**: Eliminate candidates in L_{k+1} that are infrequent, leaving only those that are frequent $\Rightarrow F_{k+1}$

3/8/2021

Introduction to Data Mining, 2nd Edition

21

21

Candidate Generation: Brute-force method

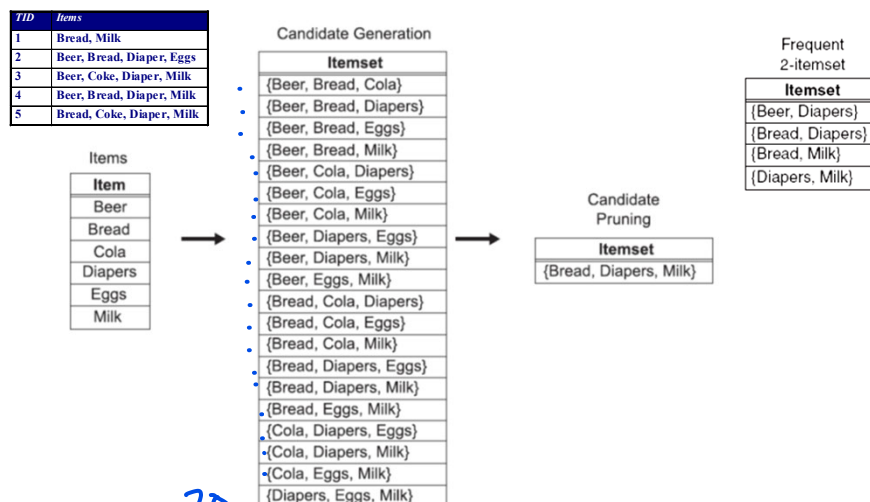


Figure 5.6. A brute-force method for generating candidate 3-itemsets.

3/8/2021

Introduction to Data Mining, 2nd Edition

22

22

Candidate Generation: Merge F_{k-1} and F₁ itemsets

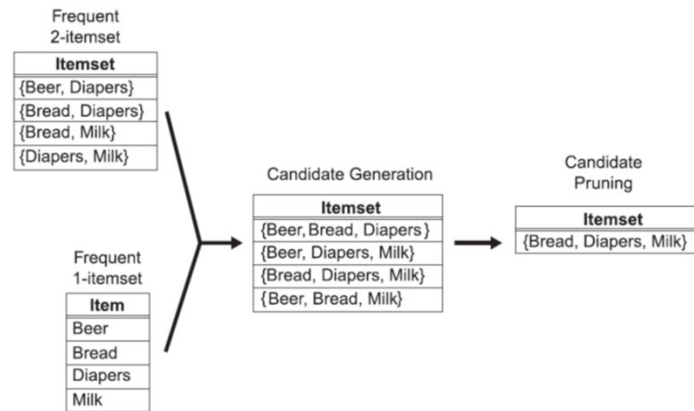


Figure 5.7. Generating and pruning candidate k -itemsets by merging a frequent $(k - 1)$ -itemset with a frequent item. Note that some of the candidates are unnecessary because their subsets are infrequent.

3/8/2021

Introduction to Data Mining, 2nd Edition

23

23

Candidate Generation: $F_{k-1} \times F_{k-1}$ Method

- Merge two frequent $(k-1)$ -itemsets if their first $(k-2)$ items are identical
- $F_3 = \{ABC, ABD, ABE, ACD, BCD, BDE, CDE\}$
 - Merge(ABC, ABD) = ABCD
 - Merge(ABC, ABE) = ABCE
 - Merge(ABD, ABE) = ABDE
 - Do not merge(ABD, ACD) because they share only prefix of length 1 instead of length 2

3/8/2021

Introduction to Data Mining, 2nd Edition

24

24

Candidate Pruning

- Let $F_3 = \{ABC, ABD, ABE, ACD, BCD, BDE, CDE\}$ be the set of frequent 3-itemsets
- $L_4 = \{ABCD, ABCE, ABDE\}$ is the set of candidate 4-itemsets generated (from previous slide)
- Candidate pruning
 - Prune ABCE because ACE and BCE are infrequent
 - Prune ABDE because ADE is infrequent
- After candidate pruning: $L_4 = \{ABCD\}$

3/8/2021

Introduction to Data Mining, 2nd Edition

25

25

Candidate Generation: Fk-1 x Fk-1 Method

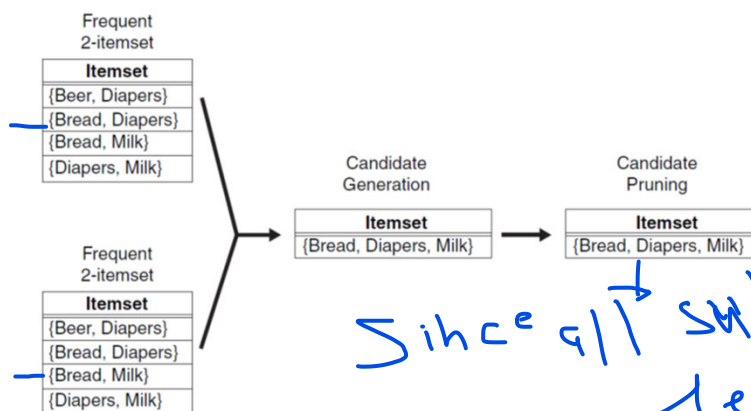


Figure 5.8. Generating and pruning candidate k -itemsets by merging pairs of frequent $(k-1)$ -itemsets.

3/8/2021

Introduction to Data Mining, 2nd Edition

26

26

using
 $F_{k-1} \times F_{k-1}$

Illustrating Apriori Principle

Item	Count
Bread	4
Coke	2
Milk	4
Beer	3
Diaper	4
Eggs	1

Items (1-itemsets)



Itemset	Count
{Bread, Milk}	3
{Bread, Beer}	2
{Bread, Diaper}	3
{Milk, Beer}	2
{Milk, Diaper}	3
{Beer, Diaper}	3

Pairs (2-itemsets)

(No need to generate candidates involving Coke or Eggs)



Itemset	Count
{Bread, Diaper, Milk}	2

Triplets (3-itemsets)

Minimum Support = 3

If every subset is considered,
 ${}^6C_1 + {}^6C_2 + {}^6C_3$
 $6 + 15 + 20 = 41$
 With support-based pruning,
 $6 + 6 + 1 = 13$

Use of $F_{k-1} \times F_{k-1}$ method for candidate generation results in only one 3-itemset. This is eliminated after the support counting step.

Alternate $F_{k-1} \times F_{k-1}$ Method

- Merge two frequent (k-1)-itemsets if the last (k-2) items of the first one is identical to the first (k-2) items of the second.

- $F_3 = \{ABC, ABD, ABE, ACD, BCD, BDE, CDE\}$
 - Merge(ABC, BCD) = ABCD
 - Merge(ABD, BDE) = ABDE
 - Merge(ACD, CDE) = ACDE
 - Merge(BCD, CDE) = BCDE

Candidate Pruning for Alternate $F_{k-1} \times F_{k-1}$ Method

- Let $F_3 = \{ABC, ABD, ABE, ACD, BCD, BDE, CDE\}$ be the set of frequent 3-itemsets
- $L_4 = \{ABCD, ABDE, ACDE, BCDE\}$ is the set of candidate 4-itemsets generated (from previous slide)
- Candidate pruning
 - Prune ABDE because ADE is infrequent
 - Prune ACDE because ACE and ADE are infrequent
 - Prune BCDE because BCE
- After candidate pruning: $L_4 = \{ABCD\}$

3/8/2021

Introduction to Data Mining, 2nd Edition

29

29

Support Counting of Candidate Itemsets

- Scan the database of transactions to determine the support of each candidate itemset
 - Must match every candidate itemset against every transaction, which is an expensive operation

TID	Items
1	Bread, Milk
2	Beer, Bread, Diaper, Eggs
3	Beer, Coke, Diaper, Milk
4	Beer, Bread, Diaper, Milk
5	Bread, Coke, Diaper, Milk

Itemset
{ Beer, Diaper, Milk}
{ Beer, Bread, Diaper}
{ Bread, Diaper, Milk}
{ Beer, Bread, Milk}

3/8/2021


Introduction to Data Mining, 2nd Edition

30

30

Support Counting of Candidate Itemsets

- To reduce number of comparisons, store the candidate itemsets in a hash structure
 - Instead of matching each transaction against every candidate, match it against candidates contained in the hashed buckets



3/8/2021

Introduction to Data Mining, 2nd Edition

31


31

Support Counting: An Example

Suppose you have 15 candidate itemsets of length 3:

{1 4 5}, {1 2 4}, {4 5 7}, {1 2 5}, {4 5 8}, {1 5 9}, {1 3 6}, {2 3 4}, {5 6 7}, {3 4 5}, {3 5 6}, {3 5 7}, {6 8 9}, {3 6 7}, {3 6 8}

How many of these itemsets are supported by transaction (1,2,3,5,6)?



3/8/2021

Introduction to Data Mining, 2nd Edition

32

32



Support Counting Using a Hash Tree

Suppose you have 15 candidate itemsets of length 3.

{1 4 5}, {1 2 4}, {4 5 7}, {1 2 5}, {4 5 8}, {1 5 9}, {1 3 6}, {2 3 4}, {5 6 7}, {3 4 5},
{3 5 6}, {3 5 7}, {6 8 9}, {3 5 7}, {3 6 8}

You need:

- Hash function
- Max leaf size: max number of itemsets stored in a leaf node (if number of candidate itemsets exceeds max leaf size, split the node)



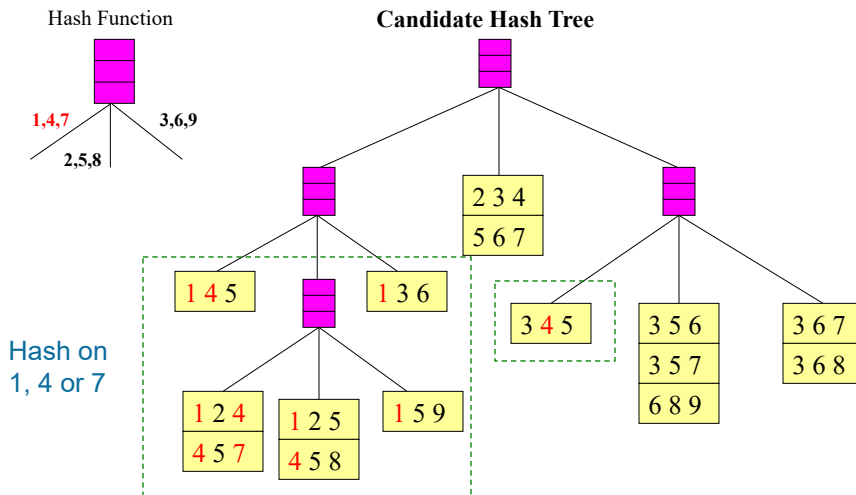
3/8/2021

Introduction to Data Mining, 2nd Edition

33

33

Support Counting Using a Hash Tree



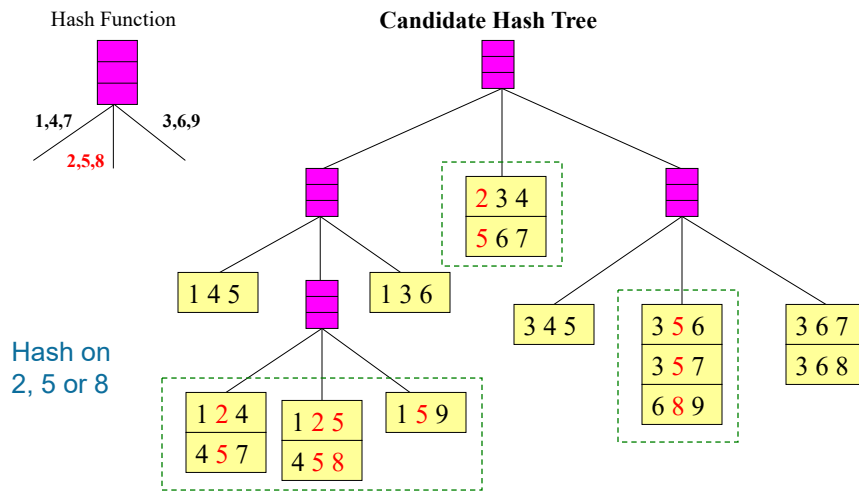
3/8/2021

Introduction to Data Mining, 2nd Edition

34

34

Support Counting Using a Hash Tree




3/8/2021

Introduction to Data Mining, 2nd Edition

35

35

Support Counting Using a Hash Tree



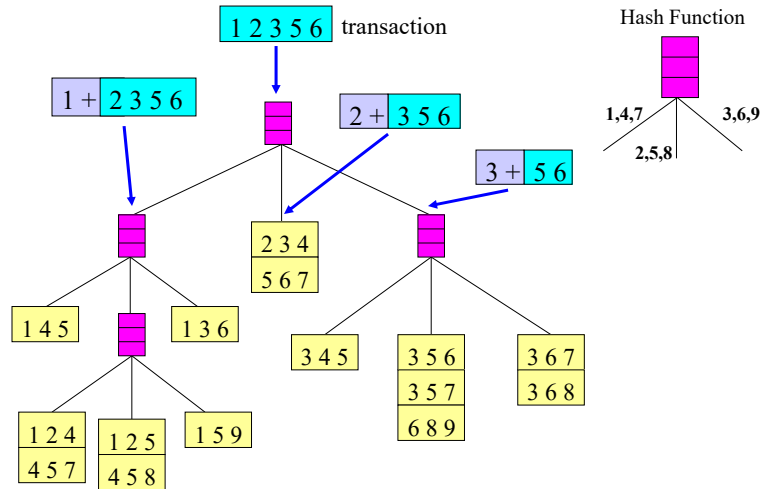
3/8/2021

Introduction to Data Mining, 2nd Edition

36

36

Support Counting Using a Hash Tree



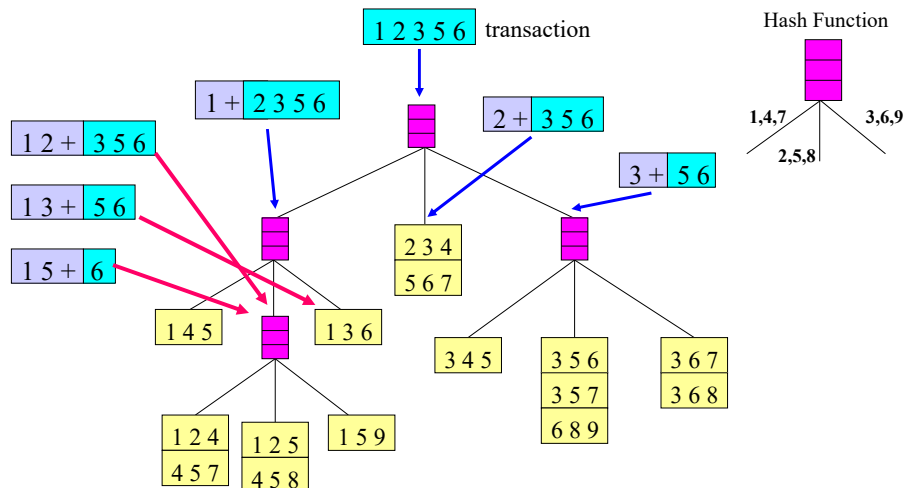
3/8/2021

Introduction to Data Mining, 2nd Edition

37

37

Support Counting Using a Hash Tree



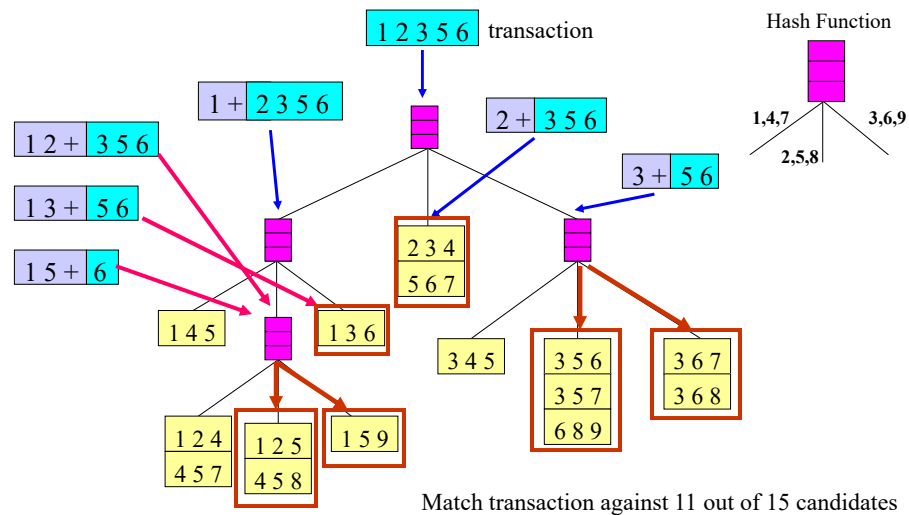
3/8/2021

Introduction to Data Mining, 2nd Edition

38

38

Support Counting Using a Hash Tree



3/8/2021

Introduction to Data Mining, 2nd Edition

39

39

Rule Generation

- Given a frequent itemset L , find all non-empty subsets $f \subset L$ such that $f \rightarrow L - f$ satisfies the minimum confidence requirement
 - If $\{A, B, C, D\}$ is a frequent itemset, candidate rules:

$ABC \rightarrow D,$	$ABD \rightarrow C,$	$ACD \rightarrow B,$	$BCD \rightarrow A,$
$A \rightarrow BCD,$	$B \rightarrow ACD,$	$C \rightarrow ABD,$	$D \rightarrow ABC$
$AB \rightarrow CD,$	$AC \rightarrow BD,$	$AD \rightarrow BC,$	$BC \rightarrow AD,$
$BD \rightarrow AC,$	$CD \rightarrow AB,$		
- If $|L| = k$, then there are $2^k - 2$ candidate association rules (ignoring $L \rightarrow \emptyset$ and $\emptyset \rightarrow L$)

3/8/2021

Introduction to Data Mining, 2nd Edition

40

40

Rule Generation

- In general, confidence does not have an anti-monotone property

$c(ABC \rightarrow D)$ can be larger or smaller than $c(AB \rightarrow D)$

- But confidence of rules generated from the same itemset has an anti-monotone property

– E.g., Suppose $\{A, B, C, D\}$ is a frequent 4-itemset:

$$c(ABC \rightarrow D) \geq c(AB \rightarrow CD) \geq c(A \rightarrow BCD)$$

– Confidence is anti-monotone w.r.t. number of items on the RHS of the rule

3/8/2021

Introduction to Data Mining, 2nd Edition


41

41

Rule Generation for Apriori Algorithm

Lattice of rules

Low
Confidence
Rule



3/8/2021

Introduction to Data Mining, 2nd Edition

42

42

Association Analysis: Basic Concepts and Algorithms

Algorithms and Complexity

43

Factors Affecting Complexity of Apriori

- Choice of minimum support threshold
- Dimensionality (number of items) of the data set
- Size of database
- Average transaction width

44

Factors Affecting Complexity of Apriori

- Choice of minimum support threshold
 - lowering support threshold results in more frequent itemsets
 - this may increase number of candidates and max length of frequent itemsets
- Dimensionality (number of items) of the data set
 -
- Size of database
 -
- Average transaction width
 -

TID	Items
1	Bread, Milk
2	Beer, Bread, Diaper, Eggs
3	Beer, Coke, Diaper, Milk
4	Beer, Bread, Diaper, Milk
5	Bread, Coke, Diaper, Milk

3/8/2021

Introduction to Data Mining, 2nd Edition

45

45

Impact of Support Based Pruning

TID	Items
1	Bread, Milk
2	Beer, Bread, Diaper, Eggs
3	Beer, Coke, Diaper, Milk
4	Beer, Bread, Diaper, Milk
5	Bread, Coke, Diaper, Milk



Items (1-itemsets)

Item	Count
Bread	4
Coke	2
Milk	4
Beer	3
Diaper	4
Eggs	1

Minimum Support = 3

If every subset is considered,
 ${}^6C_1 + {}^6C_2 + {}^6C_3$
 $6 + 15 + 20 = 41$
 With support-based pruning,
 $6 + 6 + 4 = 16$

Minimum Support = 2

If every subset is considered,
 ${}^6C_1 + {}^6C_2 + {}^6C_3 + {}^6C_4$
 $6 + 15 + 20 + 15 = 56$

3/8/2021

Introduction to Data Mining, 2nd Edition

46

46

Factors Affecting Complexity of Apriori

- Choice of minimum support threshold
 - lowering support threshold results in more frequent itemsets
 - this may increase number of candidates and max length of frequent itemsets
- Dimensionality (number of items) of the data set
 - More space is needed to store support count of itemsets
 - if number of frequent itemsets also increases, both computation and I/O costs may also increase
- Size of database
- Average transaction width
 -

TID	Items
1	Bread, Milk
2	Beer, Bread, Diaper, Eggs
3	Beer, Coke, Diaper, Milk
4	Beer, Bread, Diaper, Milk
5	Bread, Coke, Diaper, Milk

3/8/2021

Introduction to Data Mining, 2nd Edition

47

47

Factors Affecting Complexity of Apriori

- Choice of minimum support threshold
 - lowering support threshold results in more frequent itemsets
 - this may increase number of candidates and max length of frequent itemsets
- Dimensionality (number of items) of the data set
 - More space is needed to store support count of itemsets
 - if number of frequent itemsets also increases, both computation and I/O costs may also increase
- Size of database
 - run time of algorithm increases with number of transactions
- Average transaction width

TID	Items
1	Bread, Milk
2	Beer, Bread, Diaper, Eggs
3	Beer, Coke, Diaper, Milk
4	Beer, Bread, Diaper, Milk
5	Bread, Coke, Diaper, Milk

3/8/2021

Introduction to Data Mining, 2nd Edition

48

Factors Affecting Complexity of Apriori

- Choice of minimum support threshold
 - lowering support threshold results in more frequent itemsets
 - this may increase number of candidates and max length of frequent itemsets
- Dimensionality (number of items) of the data set
 - More space is needed to store support count of itemsets
 - if number of frequent itemsets also increases, both computation and I/O costs may also increase
- Size of database
 - run time of algorithm increases with number of transactions
- Average transaction width
 - transaction width increases the max length of frequent itemsets
 - number of subsets in a transaction increases with its width, increasing computation time for support counting


3/8/2021

Introduction to Data Mining, 2nd Edition


49

49

Factors Affecting Complexity of Apriori




(a) Number of candidate itemsets.




(b) Number of frequent itemsets.

Figure 6.13. Effect of support threshold on the number of candidate and frequent itemsets.



(a) Number of candidate itemsets.



(b) Number of frequent itemsets.

Figure 6.14. Effect of average transaction width on the number of candidate and frequent itemsets.

3/8/2021

Introduction to Data Mining, 2nd Edition

50

50

Compact Representation of Frequent Itemsets

- Some frequent itemsets are redundant because their supersets are also frequent

Consider the following data set. Assume support threshold = 5

TID	A1	A2	A3	A4	A5	A6	A7	A8	A9	A10	B1	B2	B3	B4	B5	B6	B7	B8	B9	B10	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
1	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
2	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
3	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
4	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
5	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
6	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0
7	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0
8	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0
9	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0
10	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0
11	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1
12	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1
13	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1
14	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1
15	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1

$$\text{Number of frequent itemsets} = 3 \times \sum_{k=1}^{10} \binom{10}{k}$$

- Need a compact representation

3/8/2021

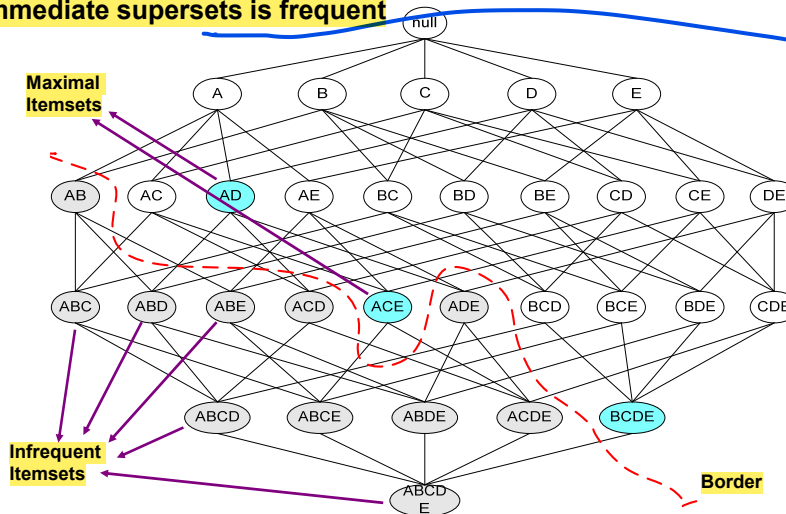
Introduction to Data Mining, 2nd Edition

51

51

Maximal Frequent Itemset

An itemset is maximal frequent if it is frequent and none of its immediate supersets is frequent



3/8/2021

Introduction to Data Mining, 2nd Edition

52

52



What are the Maximal Frequent Itemsets in this Data?

TID	A1	A2	A3	A4	A5	A6	A7	A8	A9	A10	B1	B2	B3	B4	B5	B6	B7	B8	B9	B10	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
1	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
2	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
3	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
4	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
5	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
6	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0
7	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0
8	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0
9	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0
10	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0
11	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1
12	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1
13	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1
14	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1
15	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1

Minimum support threshold = 5

(A1-A10)
(B1-B10)
(C1-C10)

3/8/2021

Introduction to Data Mining, 2nd Edition

53

53

An illustrative example

		Items									
		A	B	C	D	E	F	G	H	I	J
Transactions	1										
	2										
	3										
	4										
	5										
	6										
	7										
	8										
	9										
	10										

Support threshold (by count) : 5

Frequent itemsets: ?
Maximal itemsets: ?

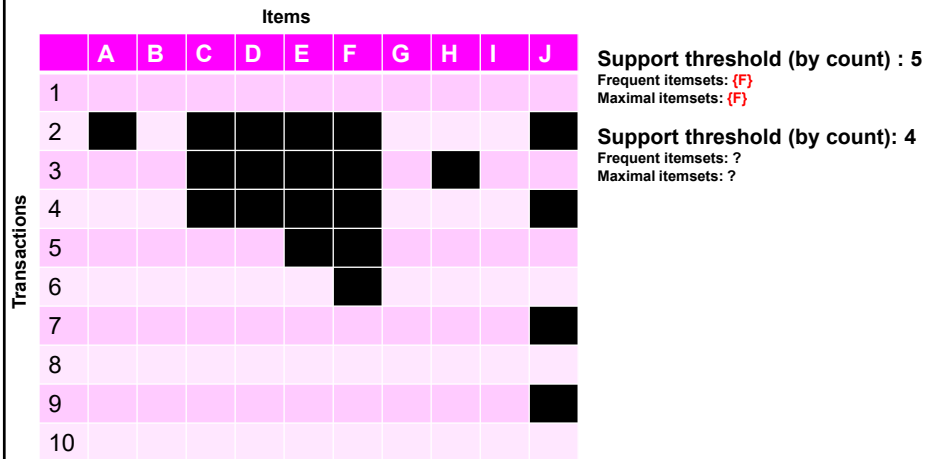
3/8/2021

Introduction to Data Mining, 2nd Edition

54

54

An illustrative example



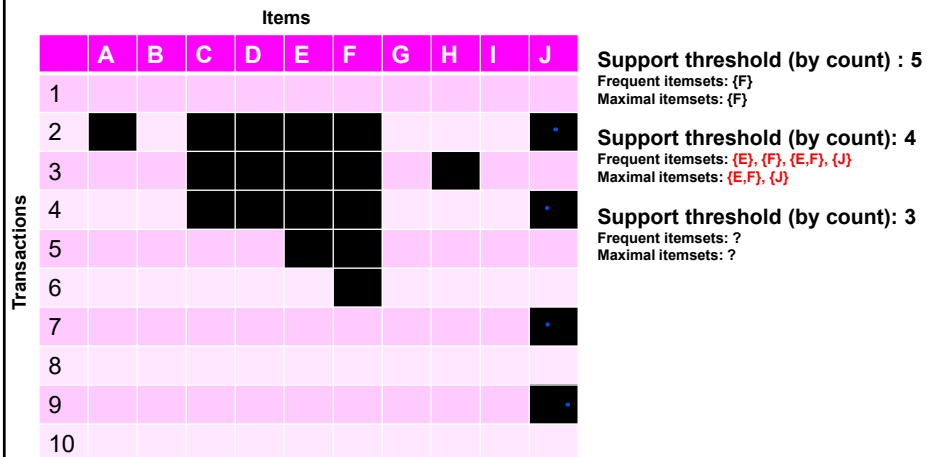
3/8/2021

Introduction to Data Mining, 2nd Edition

55

55

An illustrative example



3/8/2021

Introduction to Data Mining, 2nd Edition

56

56

An illustrative example

		Items									
Transactions		A	B	C	D	E	F	G	H	I	J
	1										
	2										
	3										
	4										
	5										
	6										
	7										
	8										
	9										
	10										

Support threshold (by count) : 5
 Frequent itemsets: {F}
 Maximal itemsets: {F}

Support threshold (by count): 4
 Frequent itemsets: {E}, {F}, {E,F}, {J}
 Maximal itemsets: {E,F}, {J}

Support threshold (by count): 3
 Frequent itemsets:
 All subsets of {C,D,E,F} + {J}
 Maximal itemsets:
 {C,D,E,F}, {J}

Another illustrative example

		Items									
Transactions		A	B	C	D	E	F	G	H	I	J
	1										
	2										
	3										
	4										
	5										
	6										
	7										
	8										
	9										
	10										

Support threshold (by count) : 5
 Maximal itemsets: {A}, {B}, {C}

Support threshold (by count): 4
 Maximal itemsets: {A,B}, {A,C}, {B,C}

Support threshold (by count): 3
 Maximal itemsets: {A,B,C}

Closed Itemset

- An itemset X is closed if none of its immediate supersets has the same support as the itemset X.
- X is not closed if at least one of its immediate supersets has support count as X.

Closed Itemset

- An itemset X is closed if none of its immediate supersets has the same support as the itemset X.
- X is not closed if at least one of its immediate supersets has support count as X.

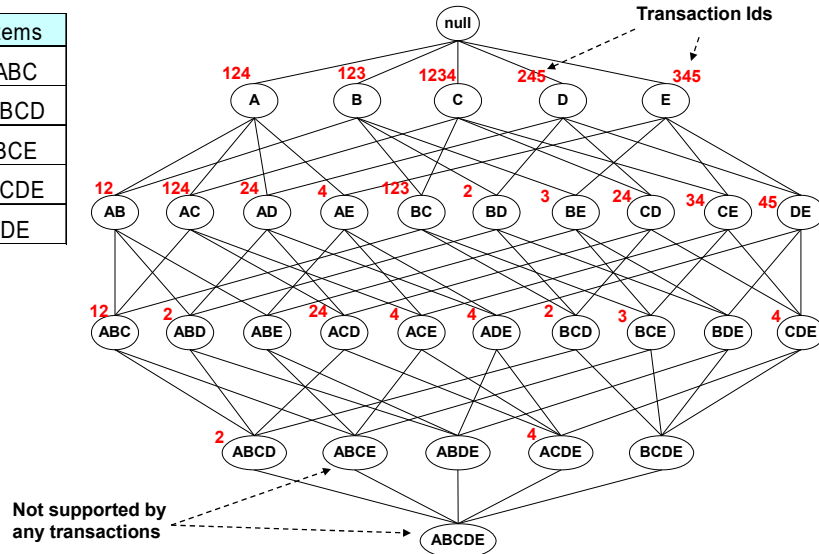
TID	Items
1	{A,B}
2	{B,C,D}
3	{A,B,C,D}
4	{A,B,D}
5	{A,B,C,D}

Itemset	Support
{A}	4
{B}	5
{C}	3
{D}	4
{A,B}	4
{A,C}	2
{A,D}	3
{B,C}	3
{B,D}	4
{C,D}	3

Itemset	Support
{A,B,C}	2
{A,B,D}	3
{A,C,D}	2
{B,C,D}	2
{A,B,C,D}	2

Maximal vs Closed Itemsets

TID	Items
1	ABC
2	ABCD
3	BCE
4	ACDE
5	DE



3/8/2021


Introduction to Data Mining, 2nd Edition

61

61

Maximal Frequent vs Closed Frequent Itemsets

TID	Items
1	ABC
2	ABCD
3	BCE
4	ACDE
5	DE



3/8/2021

Introduction to Data Mining, 2nd Edition

62

62

What are the Closed Itemsets in this Data?

TID	A1	A2	A3	A4	A5	A6	A7	A8	A9	A10	B1	B2	B3	B4	B5	B6	B7	B8	B9	B10	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10
1	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
2	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
3	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
4	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
5	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
6	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0
7	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0
8	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0
9	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0
10	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0
11	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1
12	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1
13	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1
14	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1
15	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1	1

(A1-A10)
(B1-B10)
(C1-C10)

3/8/2021

Introduction to Data Mining, 2nd Edition

63

63

Example 1

		Items												
Transactions		A	B	C	D	E	F	G	H	I	J	Itemsets	Support (counts)	Closed itemsets
	1													
	2											{C}	3	
	3											{D}	2	
	4											{C,D}	2	
	5													
	6													
	7													
	8													
	9													
	10													

3/8/2021

Introduction to Data Mining, 2nd Edition

64

64

Example 1

		Items										Transactions	
		A	B	C	D	E	F	G	H	I	J		
1													
2													
3													
4													
5													
6													
7													
8													
9													
10													

Itemsets	Support (counts)	Closed itemsets
{C}	3	✓
{D}	2	
{C,D}	2	✓

3/8/2021

Introduction to Data Mining, 2nd Edition

65

65

Example 2

		Items										Transactions	
		A	B	C	D	E	F	G	H	I	J		
1													
2													
3													
4													
5													
6													
7													
8													
9													
10													

Itemsets	Support (counts)	Closed itemsets
{C}	3	
{D}	2	
{E}	2	
{C,D}	2	
{C,E}	2	
{D,E}	2	
{C,D,E}	2	

3/8/2021

Introduction to Data Mining, 2nd Edition

66

66

Example 2

		Items									
Transactions		A	B	C	D	E	F	G	H	I	J
	1										
	2										
	3										
	4										
	5										
	6										
	7										
	8										
	9										
	10										

Itemsets	Support (counts)	Closed itemsets
{C}	3	✓
{D}	2	
{E}	2	
{C,D}	2	
{C,E}	2	
{D,E}	2	
{C,D,E}	2	✓

3/8/2021

Introduction to Data Mining, 2nd Edition

67

67

Example 3

		Items									
Transactions		A	B	C	D	E	F	G	H	I	J
	1										
	2										
	3										
	4										
	5										
	6										
	7										
	8										
	9										
	10										

Closed itemsets: {C,D,E,F}, {C,F}

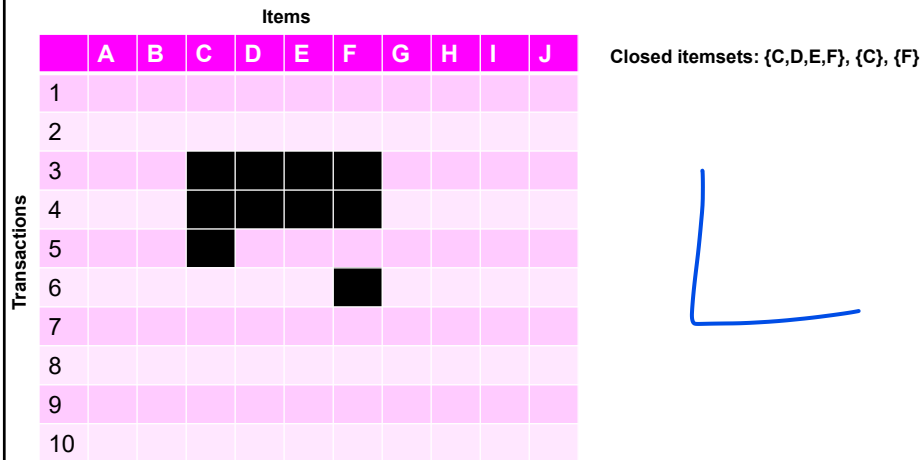
3/8/2021

Introduction to Data Mining, 2nd Edition

68

68

Example 4



3/8/2021

Introduction to Data Mining, 2nd Edition

69

69

Maximal vs Closed Itemsets

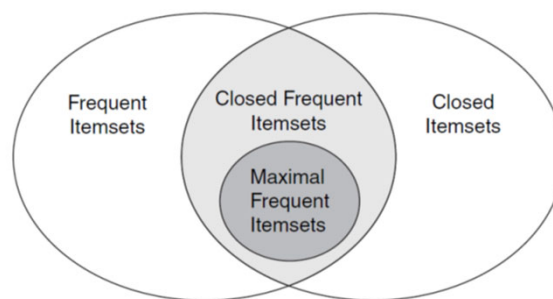


Figure 5.18. Relationships among frequent, closed, closed frequent, and maximal frequent itemsets.

3/8/2021


Introduction to Data Mining, 2nd Edition

70

70

Example question

- Given the following transaction data sets (dark cells indicate presence of an item in a transaction) and a support threshold of 20%, answer the following questions



- What is the number of frequent itemsets for each dataset? Which dataset will produce the most number of frequent itemsets?
- Which dataset will produce the longest frequent itemset?
- Which dataset will produce frequent itemsets with highest maximum support?
- Which dataset will produce frequent itemsets containing items with widely varying support levels (i.e., itemsets containing items with mixed support, ranging from 20% to more than 70%)?
- What is the number of maximal frequent itemsets for each dataset? Which dataset will produce the most number of maximal frequent itemsets?
- What is the number of closed frequent itemsets for each dataset? Which dataset will produce the most number of closed frequent itemsets?

3/8/2021

Introduction to Data Mining, 2nd Edition

71

71

Pattern Evaluation

- Association rule algorithms can produce large number of rules
- Interestingness measures can be used to prune/rank the patterns
 - In the original formulation, support & confidence are the only measures used

3/8/2021

Introduction to Data Mining, 2nd Edition

72

72

Computing Interestingness Measure

- Given $X \rightarrow Y$ or $\{X, Y\}$, information needed to compute interestingness can be obtained from a contingency table

Contingency table

	Y	\bar{Y}	
X	f_{11}	f_{10}	f_{1+}
\bar{X}	f_{01}	f_{00}	f_{0+}
	f_{+1}	f_{+0}	N

f_{11} : support of X and Y

f_{10} : support of \bar{X} and \bar{Y}

f_{01} : support of \bar{X} and Y

f_{00} : support of X and \bar{Y}

Used to define various measures

- support, confidence, Gini, entropy, etc.

Drawback of Confidence

Custo mers	Tea	Coffee	...
C1	0	1	...
C2	1	0	...
C3	1	1	...
C4	1	0	...
...			

	Coffee	\overline{Coffee}	
Tea	150	50	200
\overline{Tea}	650	150	800
	800	200	1000

Association Rule: Tea \rightarrow Coffee

Confidence $\cong P(\text{Coffee}|\text{Tea}) = 150/200 = 0.75$

Confidence > 50%, meaning people who drink tea are more likely to drink coffee than not drink coffee

So rule seems reasonable

Drawback of Confidence

	Coffee	$\overline{\text{Coffee}}$	
Tea	150	50	200
$\overline{\text{Tea}}$	650	150	800
	800	200	1000

Association Rule: Tea \rightarrow Coffee

Confidence = $P(\text{Coffee}|\text{Tea}) = 150/200 = 0.75$

but $P(\text{Coffee}) = 0.8$, which means knowing that a person drinks tea reduces the probability that the person drinks coffee!

\Rightarrow Note that $P(\text{Coffee}|\overline{\text{Tea}}) = 650/800 = 0.8125$

Drawback of Confidence

Custo mers	Tea	Honey	...
C1	0	1	...
C2	1	0	...
C3	1	1	...
C4	1	0	...
...			

	Honey	$\overline{\text{Honey}}$	
Tea	100	100	200
$\overline{\text{Tea}}$	20	780	800
	120	880	1000

Association Rule: Tea \rightarrow Honey

Confidence $\cong P(\text{Honey}|\text{Tea}) = 100/200 = 0.50$

Confidence = 50%, which may mean that drinking tea has little influence whether honey is used or not

So rule seems uninteresting

But $P(\text{Honey}) = 120/1000 = .12$ (hence tea drinkers are far more likely to have honey)

Measure for Association Rules

- So, what kind of rules do we really want?
 - Confidence($X \rightarrow Y$) should be sufficiently high
 - ◆ To ensure that people who buy X will more likely buy Y than not buy Y
 - Confidence($X \rightarrow Y$) > support(Y)
 - ◆ Otherwise, rule will be misleading because having item X actually reduces the chance of having item Y in the same transaction
 - ◆ Is there any measure that capture this constraint?
 - Answer: Yes. There are many of them.

3/8/2021

Introduction to Data Mining, 2nd Edition

77

77

Statistical Relationship between X and Y

- The criterion
 $\text{confidence}(X \rightarrow Y) = \text{support}(Y)$
is equivalent to:
 - $P(Y|X) = P(Y)$
 - $P(X, Y) = P(X) \times P(Y)$ (X and Y are independent)
- If $P(X, Y) > P(X) \times P(Y)$: X & Y are positively correlated
- If $P(X, Y) < P(X) \times P(Y)$: X & Y are negatively correlated

3/8/2021

Introduction to Data Mining, 2nd Edition

78

78

Measures that take into account statistical dependence

$$Lift = \frac{P(Y|X)}{P(Y)}$$

$$Interest = \frac{P(X,Y)}{P(X)P(Y)}$$

$$PS = P(X,Y) - P(X)P(Y)$$

$$\phi\text{-coefficient} = \frac{P(X,Y) - P(X)P(Y)}{\sqrt{P(X)[1 - P(X)]P(Y)[1 - P(Y)]}}$$

lift is used for rules while
interest is used for itemsets

3/8/2021

Introduction to Data Mining, 2nd Edition

79

79

Example: Lift/Interest

	Coffee	<u>Coffee</u>	
<u>Tea</u>	150	50	200
Tea	650	150	800
	800	200	1000

Association Rule: Tea → Coffee

Confidence = P(Coffee|Tea) = 0.75

but P(Coffee) = 0.8

⇒ Interest = 0.15 / (0.2 × 0.8) = 0.9375 (< 1, therefore is negatively associated)

So, is it enough to use confidence/Interest for pruning?

3/8/2021

Introduction to Data Mining, 2nd Edition

80

80

There are lots of measures proposed in the literature

Measure (Symbol)	Definition
Correlation (ϕ)	$\frac{N f_{11} - f_{1+} f_{+1}}{\sqrt{f_{1+} f_{+1} f_{0+} f_{+0}}}$
Odds ratio (α)	$(f_{11} f_{00}) / (f_{10} f_{01})$
Kappa (κ)	$\frac{N f_{11} + N f_{00} - f_{1+} f_{+1} - f_{0+} f_{+0}}{N^2 - f_{1+} f_{+1} - f_{0+} f_{+0}}$
Interest (I)	$(N f_{11}) / (f_{1+} f_{+1})$
Cosine (IS)	$(f_{11}) / (\sqrt{f_{1+} f_{+1}})$
Piatetsky-Shapiro (PS)	$\frac{f_{11}}{N} - \frac{f_{1+} f_{+1}}{N^2}$
Collective strength (S)	$\frac{f_{11} + f_{00}}{f_{1+} f_{+1} + f_{0+} f_{+0}} \times \frac{N - f_{1+} f_{+1} - f_{0+} f_{+0}}{N - f_{11} - f_{00}}$
Jaccard (ζ)	$f_{11} / (f_{1+} + f_{+1} - f_{11})$
All-confidence (h)	$\min \left[\frac{f_{11}}{f_{1+}}, \frac{f_{11}}{f_{+1}} \right]$

Comparing Different Measures

10 examples of contingency tables:

Example	f_{11}	f_{10}	f_{01}	f_{00}
E1	8123	83	424	1370
E2	8330	2	622	1046
E3	9481	94	127	298
E4	3954	3080	5	2961
E5	2886	1363	1320	4431
E6	1500	2000	500	6000
E7	4000	2000	1000	3000
E8	4000	2000	2000	2000
E9	1720	7121	5	1154
E10	61	2483	4	7452

Rankings of contingency tables using various measures:

	ϕ	α	κ	I	IS	PS	S	ζ	h
E_1	1	3	1	6	2	2	1	2	2
E_2	2	1	2	7	3	5	2	3	3
E_3	3	2	4	4	5	1	3	6	8
E_4	4	8	3	3	7	3	4	7	5
E_5	5	7	6	2	9	6	6	9	9
E_6	6	9	5	5	6	4	5	5	7
E_7	7	6	7	9	1	8	7	1	1
E_8	8	10	8	8	8	7	8	8	7
E_9	9	4	9	10	4	9	9	4	4
E_{10}	10	5	10	1	10	10	10	10	10

Property under Inversion Operation

	A	B	\bar{A}	\bar{B}
Transaction 1 →	1	0	0	1
■	0	0	1	1
■	0	0	1	1
■	0	0	1	1
■	0	1	1	0
■	0	0	1	1
■	0	0	1	1
■	0	0	1	1
Transaction N →	1	0	0	1

(a) (b)

3/8/2021

Introduction to Data Mining, 2nd Edition

83

83

Property under Inversion Operation

	A	B	\bar{A}	\bar{B}
Transaction 1 →	1	0	0	1
■	0	0	1	1
■	0	0	1	1
■	0	0	1	1
■	0	1	1	0
■	0	0	1	1
■	0	0	1	1
■	0	0	1	1
Transaction N →	1	0	0	1

(a) (b)

Correlation:
IS/cosine

-0.1667
0.0

-0.1667
0.825

3/8/2021


Introduction to Data Mining, 2nd Edition

84

84

Property under Null Addition

	B	\overline{B}	
A	700	100	800
\overline{A}	100	100	200
	800	200	1000



	B	\overline{B}	
A	700	100	800
\overline{A}	10	1100	1200
	800	1200	2000

Invariant measures:

- ◆ cosine, Jaccard, All-confidence, confidence

Non-invariant measures:


- ◆ correlation, Interest/Lift, odds ratio, etc


Property under Row/Column Scaling

Grade-Gender Example (Mosteller, 1968):

	Male	Female	
High	30	20	50
Low	40	10	50
	70	30	100

	Male	Female	
High	60	60	120
Low	80	30	110
	140	90	230


 2x


 3x

Mosteller:

Underlying association should be independent of the relative number of male and female students in the samples

Odds-Ratio $((f_{11}+f_{00})/(f_{10}+f_{01}))$ has this property

Property under Row/Column Scaling

Relationship between Mask use and susceptibility to Covid:

	Covid-Positive	Covid-Free			Covid-Positive	Covid-Free	
Mask	20	30	50	Mask	40	300	340
No-Mask	40	10	50	No-Mask	80	100	180
	60	40	100		120	400	520
					2x	10x	

Mosteller:

Underlying association should be independent of the relative number of Covid-positive and Covid-free subjects

Odds-Ratio $((f_{11}f_{00})/(f_{10}f_{01}))$ has this property

3/8/2021

Introduction to Data Mining, 2nd Edition

87

87

Different Measures have Different Properties

Symbol	Measure	Inversion	Null Addition	Scaling
ϕ	ϕ -coefficient	Yes	No	No
α	odds ratio	Yes	No	Yes
κ	Cohen's	Yes	No	No
I	Interest	No	No	No
IS	Cosine	No	Yes	No
PS	Piatetsky-Shapiro's	Yes	No	No
S	Collective strength	Yes	No	No
ζ	Jaccard	No	Yes	No
h	All-confidence	No	Yes	No
s	Support	No	No	No

3/8/2021

Introduction to Data Mining, 2nd Edition

88

88

Simpson's Paradox

- Observed relationship in data may be influenced by the presence of other confounding factors (hidden variables)
 - Hidden variables may cause the observed relationship to disappear or reverse its direction!
- Proper stratification is needed to avoid generating spurious patterns

Simpson's Paradox

- Recovery rate from Covid
 - Hospital A: 80%
 - Hospital B: 90%
- Which hospital is better?

Simpson's Paradox

- Recovery rate from Covid
 - Hospital A: 80%
 - Hospital B: 90%
- Which hospital is better?
- Covid recovery rate on older population
 - Hospital A: 50%
 - Hospital B: 30%
- Covid recovery rate on younger population
 - Hospital A: 99%
 - Hospital B: 98%

Simpson's Paradox

- Covid-19 death: (per 100,000 of population)
 - County A: 15
 - County B: 10
- Which state is managing the pandemic better?

Simpson's Paradox

- Covid-19 death: (per 100,000 of population)
 - County A: 15
 - County B: 10
- Which state is managing the pandemic better?
- Covid death rate on older population
 - County A: 20
 - County B: 40
- Covid death rate on younger population
 - County A: 2
 - County B: 5

3/8/2021

Introduction to Data Mining, 2nd Edition


93

93

Effect of Support Distribution on Association Mining

- Many real data sets have skewed support distribution

Support
distribution of
a retail data set



3/8/2021

Introduction to Data Mining, 2nd Edition


94

94

Effect of Support Distribution

- Difficult to set the appropriate *minsup* threshold
 - If *minsup* is too high, we could miss itemsets involving interesting rare items (e.g., {caviar, vodka})
 - If *minsup* is too low, it is computationally expensive and the number of itemsets is very large

Cross-Support Patterns



A cross-support pattern involves items with varying degree of support

- Example: {caviar,milk}

How to avoid such patterns?

A Measure of Cross Support


- Given an itemset, $X = \{x_1, x_2, \dots, x_d\}$, with d items, we can define a measure of cross support, r , for the itemset

$$r(X) = \frac{\min\{s(x_1), s(x_2), \dots, s(x_d)\}}{\max\{s(x_1), s(x_2), \dots, s(x_d)\}}$$

where $s(x_i)$ is the support of item x_i

- Can use $r(X)$ to prune cross support patterns

Confidence and Cross-Support Patterns



Observation:

$\text{conf}(\text{caviar} \rightarrow \text{milk})$ is very high
but
 $\text{conf}(\text{milk} \rightarrow \text{caviar})$ is very low

Therefore,

$\min(\text{conf}(\text{caviar} \rightarrow \text{milk}), \text{conf}(\text{milk} \rightarrow \text{caviar}))$

is also very low

H-Confidence

- To avoid patterns whose items have very different support, define a new evaluation measure for itemsets
 - Known as h-confidence or all-confidence
- Specifically, given an itemset $X = \{x_1, x_2, \dots, x_d\}$
 - h-confidence is the minimum confidence of any association rule formed from itemset X
 - $\text{hconf}(X) = \min(\text{conf}(X_1 \rightarrow X_2))$,
 where $X_1, X_2 \subset X, X_1 \cap X_2 = \emptyset, X_1 \cup X_2 = X$
 For example: $X_1 = \{x_1, x_2\}, X_2 = \{x_3, \dots, x_d\}$

3/8/2021

Introduction to Data Mining, 2nd Edition

99

99

H-Confidence ...

- But, given an itemset $X = \{x_1, x_2, \dots, x_d\}$
 - What is the lowest confidence rule you can obtain from X ?
 - Recall $\text{conf}(X_1 \rightarrow X_2) = s(X_1 \cup X_2) / \text{support}(X_1)$
 - ◆ The numerator is fixed: $s(X_1 \cup X_2) = s(X)$
 - ◆ Thus, to find the lowest confidence rule, we need to find the X_1 with highest support
 - ◆ Consider only rules where X_1 is a single item, i.e., $\{x_1\} \rightarrow X - \{x_1\}, \{x_2\} \rightarrow X - \{x_2\}, \dots$, or $\{x_d\} \rightarrow X - \{x_d\}$

$$\begin{aligned} \text{hconf}(X) &= \min \left\{ \frac{s(X)}{s(x_1)}, \frac{s(X)}{s(x_2)}, \dots, \frac{s(X)}{s(x_d)} \right\} \\ &= \frac{s(X)}{\max\{s(x_1), s(x_2), \dots, s(x_d)\}} \end{aligned}$$

3/8/2021

Introduction to Data Mining, 2nd Edition

100

100

Cross Support and H-confidence

- By the anti-montone property of support

$$s(X) \leq \min\{s(x_1), s(x_2), \dots, s(x_d)\}$$

- Therefore, we can derive a relationship between the h-confidence and cross support of an itemset

$$\begin{aligned} \text{hconf}(X) &= \frac{s(X)}{\max\{s(x_1), s(x_2), \dots, s(x_d)\}} \\ &\leq \frac{\min\{s(x_1), s(x_2), \dots, s(x_d)\}}{\max\{s(x_1), s(x_2), \dots, s(x_d)\}} \\ &= r(X) \end{aligned}$$

Thus, $\text{hconf}(X) \leq r(X)$

Cross Support and H-confidence ...

- Since, $\text{hconf}(X) \leq r(X)$, we can eliminate cross support patterns by finding patterns with h-confidence $< h_c$, a user set threshold
- Notice that

$$0 \leq \text{hconf}(X) \leq r(X) \leq 1$$

- Any itemset satisfying a given h-confidence threshold, h_c , is called a **hyperclique**
- H-confidence can be used instead of or in conjunction with support

Properties of Hypercliques

- Hypercliques are itemsets, but not necessarily frequent itemsets
 - Good for finding low support patterns
- H-confidence is anti-monotone
- Can define closed and maximal hypercliques in terms of h-confidence
 - A hyperclique X is closed if none of its immediate supersets has the same h-confidence as X
 - A hyperclique X is maximal if $\text{hconf}(X) \leq h_c$ and none of its immediate supersets, Y , have $\text{hconf}(Y) \leq h_c$

3/8/2021

Introduction to Data Mining, 2nd Edition

103

103

Properties of Hypercliques ...

- Hypercliques have the high-affinity property
 - Think of the individual items as sparse binary vectors
 - h-confidence gives us information about their pairwise Jaccard and cosine similarity
 - ◆ Assume x_1 and x_2 are any two items in an itemset X
 - ◆ $\text{Jaccard}(x_1, x_2) \geq \text{hconf}(X)/2$
 - ◆ $\cos(x_1, x_2) \geq \text{hconf}(X)$
 - Hypercliques that have a high h-confidence consist of very similar items as measured by Jaccard and cosine
- The items in a hyperclique cannot have widely different support
 - Allows for more efficient pruning

3/8/2021

Introduction to Data Mining, 2nd Edition

104

104

Example Applications of Hypercliques

- Hypercliques are used to find strongly coherent groups of items
 - Words that occur together in documents
 - Proteins in a protein interaction network

In the figure at the right, a gene ontology hierarchy for biological process shows that the identified proteins in the hyperclique (PRE2, ..., SCL1) perform the same function and are involved in the same biological process

